

# Github Repository URL

[Link to our repo](#)

## Link to our live app

[Link to our live app](#)

## 1. Problem Statement

The Dunder Mifflin Paper Company, a well-known paper supplier situated in Scranton, Pennsylvania, is suffering a disturbing trend of excessive staff turnover, which is reducing operational efficiency and worker morale. We, the analytical team, are concerned about the company's high attrition rates, which we feel are damaging to its culture and performance. This research intends to use specific employee data, including tenure, income, work satisfaction, and departmental functions, to uncover major variables influencing employee turnover. We will create a prediction model based on detailed data analysis to put employees into three risk categories for leaving: extremely probable, moderately likely, and marginally likely. Our ultimate objective is to deliver actionable information that will assist in developing successful retention strategies to minimize turnover and improve overall employee engagement. Although the setting and data are fictitious and created for instructional reasons, the approaches used will closely resemble real-world analytical applications in human resource management, providing meaningful learning benefits.

## 2. Hypothesis Generation

We hypothesize that certain employee attributes and workplace conditions have a substantial impact on employee turnover at Dunder Mifflin Paper Company. We estimate that poorer job satisfaction, insufficient work-life balance, and longer travel lengths are highly related with increased risk of quitting the organization. Furthermore, we believe that employees with longer tenure, higher income, and recent promotions are less likely to quit, indicating that job stability and recognition are important determinants in employee retention. Our investigation aims to test these ideas by using predictive modelling approaches to categorize employees into distinct churn risk groups. By validating or disputing these assumptions, we will be able to pinpoint the most important areas for action to enhance retention methods and minimize staff turnover.

## 3. Prepare Data

### Import the libraries

In [1]:

```
import numpy as np
import pandas as pd
```

### Importing the dataset

<https://www.kaggle.com/datasets/cocolicoq4/employee-churn-at-dunder-mifflin-paper-company>

In [2]:

```
df = pd.read_csv('../data/office_churn_dataset.csv')
```

```
print(df.info())
df.head()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   EmployeeID                           1543 non-null   int64
1   Branch                               1535 non-null   object
2   Tenure                               1534 non-null   float64
3   Salary                               1534 non-null   float64
4   Department                           1543 non-null   object
5   JobSatisfaction                       1515 non-null   float64
6   WorkLifeBalance                       1515 non-null   float64
7   CommuteDistance                       1543 non-null   object
8   MaritalStatus                         1543 non-null   object
9   Education                             1543 non-null   object
10  PerformanceRating                     1536 non-null   float64
11  TrainingHours                         1352 non-null   float64
12  OverTime                             1443 non-null   object
13  NumProjects                           1444 non-null   float64
14  YearsSincePromotion                   1542 non-null   float64
15  EnvironmentSatisfaction                1515 non-null   float64
16  ChurnLikelihood                       1543 non-null   object
dtypes: float64(9), int64(1), object(7)
memory usage: 205.1+ KB
None
```

Out[2]:

	EmployeeID	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus
0	1	San Francisco	4.0	63000.0	Legal	3.0	3.0	Long	Married
1	2	Chicago	14.0	72000.0	Accounting	4.0	4.0	Short	Single
2	3	Miami	4.0	40000.0	Quality Assurance	3.0	3.0	Medium	Single
3	4	Scranton	2.0	55000.0	Legal	3.0	3.5	Short	Married
4	5	Scranton	10.0	55500.0	Legal	3.0	3.0	Medium	Married

## Understanding each feature

### Dataset Features:

- **EmployeeID:** A unique identifier for each employee.
- **Branch:** The "Branch" feature represents the geographic location of each employee within one of the 12 Dunder Mifflin branches across the United States.
- **Tenure:** The number of years the employee has been with the company.
- **Salary:** The employee's annual salary.
- **Department:** The department in which the employee works (e.g., Sales, Accounting, Customer Service).
- **JobSatisfaction:** The employee's self-reported job satisfaction level (on a scale from 1 to 5, with 5 being highly satisfied).
- **WorkLifeBalance:** The employee's self-reported work-life balance rating (on a scale from 1 to 5, with 5 being excellent).
- **CommuteDistance:** The distance the employee commutes to work (e.g., Short, Medium, Long).
- **MaritalStatus:** The marital status of the employee (e.g., Single, Married, Divorced).
- **Education:** The highest level of education attained by the employee (e.g., High School, Bachelor's, Master's).
- **PerformanceRating:** The employee's performance rating (on a scale from 1 to 5, with 5 being excellent).
- **TrainingHours:** The number of hours of training the employee has received.
- **OverTime:** Whether the employee works overtime or not

- **OverTime:** Whether the employee works overtime or not.
- **NumProjects:** The number of projects the employee is currently working on.
- **YearsSincePromotion:** The number of years since the employee's last promotion.
- **EnvironmentSatisfaction:** The employee's self-reported environment satisfaction (on a scale from 1 to 5, with 5 being highly satisfied).

**Classes (Target Variable):** Employees will be classified into four classes based on their likelihood to leave the company:

- **Class A:** Highly likely to leave.
- **Class B:** Moderately likely to leave.
- **Class C:** Slightly likely to leave.

## Remove irrelevant columns

In [3]:

```

irrelevant_features = ['EmployeeID']

df.drop(columns=irrelevant_features, inplace=True)
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Branch                                1535 non-null   object
1   Tenure                                1534 non-null   float64
2   Salary                                1534 non-null   float64
3   Department                            1543 non-null   object
4   JobSatisfaction                        1515 non-null   float64
5   WorkLifeBalance                        1515 non-null   float64
6   CommuteDistance                        1543 non-null   object
7   MaritalStatus                          1543 non-null   object
8   Education                              1543 non-null   object
9   PerformanceRating                      1536 non-null   float64
10  TrainingHours                          1352 non-null   float64
11  OverTime                               1443 non-null   object
12  NumProjects                            1444 non-null   float64
13  YearsSincePromotion                    1542 non-null   float64
14  EnvironmentSatisfaction                 1515 non-null   float64
15  ChurnLikelihood                        1543 non-null   object
dtypes: float64(9), object(7)
memory usage: 193.0+ KB

```

## 4. Exploratory Data Analysis (EDA)

### i) Univariate Analysis

#### Reading random records

In [4]:

```
df.sample(5)
```

Out[4]:

	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus	Education
707	Chicago	10.0	54500.0	Quality Assurance	3.0	3.000000	Short	Married	Hig School
1155	San Francisco	3.0	73500.0	Product Management	5.0	5.000000	Short	Divorced	Bachelor

	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus	Education
638	Miami	8.0	55000.0	Human Resources	1.0	NaN	Medium	Single	Hig Scho
531	Chicago	9.0	62500.0	Legal	4.0	3.609454	Long	Divorced	Bachelo
133	Denver	3.0	54000.0	Customer Service	4.0	3.000000	Medium	Single	Hig Scho

## Analysis of numerical attributes

In [5]:

```
df.describe()
```

Out[5]:

	Tenure	Salary	JobSatisfaction	WorkLifeBalance	PerformanceRating	TrainingHours	NumProjects	YearsSi
count	1534.000000	1534.000000	1515.000000	1515.000000	1536.000000	1352.000000	1444.000000	
mean	7.612125	66654.498044	3.421782	3.770770	3.493310	34.645646	3.500687	
std	4.123834	8473.622168	1.095047	0.481407	0.411208	22.970267	0.714107	
min	0.000000	40000.000000	1.000000	2.000000	1.639834	5.196002	1.505266	
25%	5.000000	61000.000000	3.000000	3.500000	3.333333	18.000000	3.000000	
50%	7.000000	66500.000000	3.000000	4.000000	3.363192	30.000000	3.618543	
75%	10.000000	72375.000000	4.000000	4.000000	3.666667	40.000000	4.000000	
max	27.000000	98000.000000	5.000000	5.000000	5.000000	96.000000	6.394718	

## Visualization of the data

### Importing Libraries for Visualization

In [6]:


```
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns
```


### Categorical Visualization

#### Branch

In [7]:

```
df["Branch"].value_counts().plot(kind="pie", autopct='%.2f', figsize=(5, 8), y='', x='');
```






## Department

In [8]:


```
df["Department"].value_counts().plot(kind="pie", autopct='% .2f', figsize=(8, 5), y='', x='');
```



## Education

In [9]:

```
df["Education"].value_counts().plot(kind="pie", autopct='% .2f', figsize=(3, 5), y='', x='');
```




## Job Satisfaction

In [10]:

```
sns.set_theme(rc={'figure.figsize': (3,3)})
sns.histplot(x="JobSatisfaction", data=df,color="red");
```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):




## Numerical Visualization

### *Salary*

In [11]:

```
sns.set_theme(rc={'figure.figsize': (5,5)})
sns.histplot(x="Salary", data=df, kde=True,color="red");
```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):



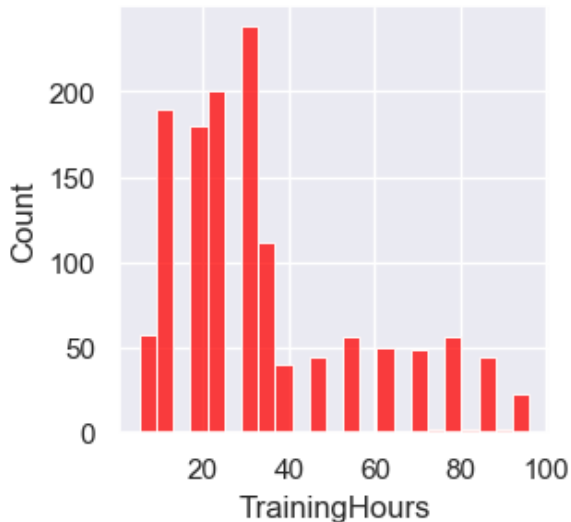
## Training Hours

In [12]:

```
sns.set_theme(rc={'figure.figsize':(3,3)})
sns.histplot(x="TrainingHours", data=df, color="red");
```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```



## Comparing Target Values for Categorical Imbalances

In [13]:

```
labels = (
    df['ChurnLikelihood']
    .astype('str')
    .str.replace('0', 'No', regex=True)
    .str.replace('1', 'Yes', regex=True)
    .value_counts()
)

fig = px.bar(
    data_frame=labels,
    x=labels.index,
    y=labels.values,
    title=f'Class Imbalance',
    color=labels.index
)

fig.update_layout(xaxis_title='Churn Status', yaxis_title='Number of Customers')
fig.show()
```

## ii) Bivariate Analysis

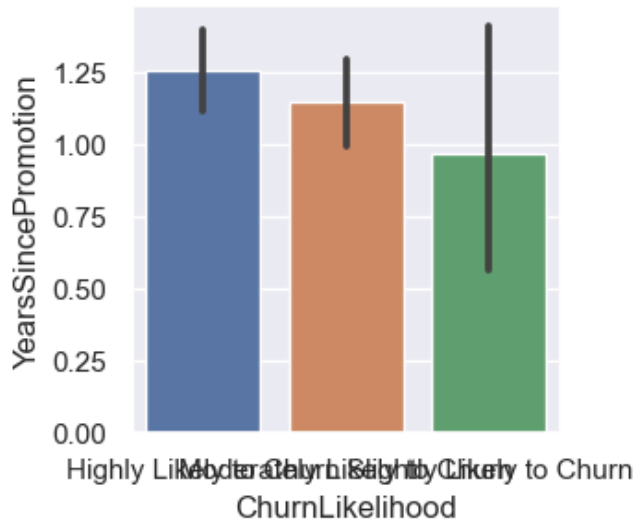
### 1. Years since promotion and Churn Likelihood

In [14]:

```
sns.barplot(x = 'ChurnLikelihood' , y = 'YearsSincePromotion' , data=df)
```

Out[14]:

<Axes: xlabel='ChurnLikelihood', ylabel='YearsSincePromotion'>




## 2 . Education and Churn likelihood

```
In [15]:
sns.heatmap(pd.crosstab(df['Education'] , df['ChurnLikelihood']))
```

Out[15]:

```
<Axes: xlabel='ChurnLikelihood', ylabel='Education'>
```



## 3. Department and Churn likelihood


```
In [16]:
sns.heatmap(pd.crosstab(df['Education'] , df['ChurnLikelihood']))
```

Out[16]:

```
<Axes: xlabel='ChurnLikelihood', ylabel='Education'>
```








#### 4. Performance rating and Salary

In [17]:

```
sns.scatterplot(x = 'PerformanceRating' , y = 'Salary' , data=df)
```

Out[17]:

<Axes: xlabel='PerformanceRating', ylabel='Salary'>



#### 5. JobSatisfaction and ChurnLikelihood


In [18]:

```
sns.barplot(x = 'JobSatisfaction' , y = 'ChurnLikelihood' , data=df)
```

Out[18]:

<Axes: xlabel='JobSatisfaction', ylabel='ChurnLikelihood'>





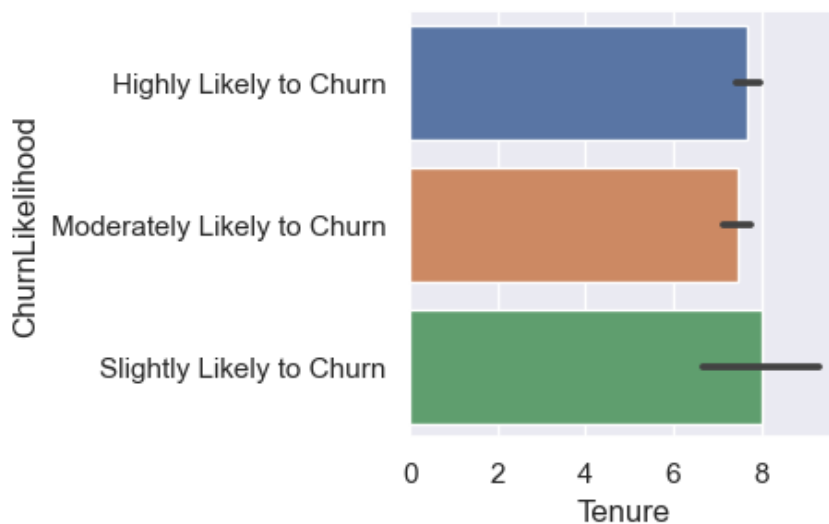
## 6. Tenure and ChurnLikelihood

In [19]:

```
sns.barplot(x = 'Tenure' , y = 'ChurnLikelihood' , data=df)
```

Out[19]:

<Axes: xlabel='Tenure', ylabel='ChurnLikelihood'>



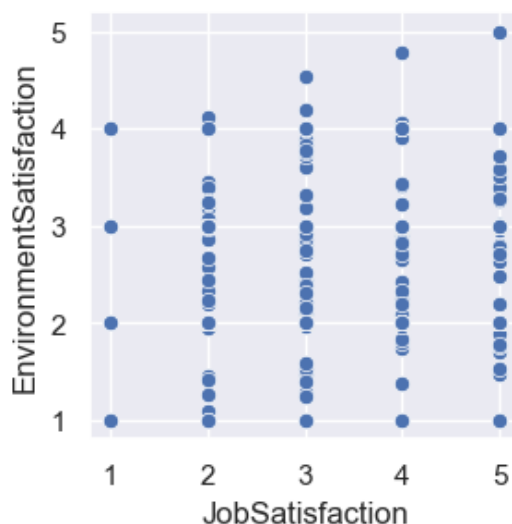
## 7. JobSatisfaction and EnvironmentSatisfaction

In [20]:

```
sns.scatterplot(x = 'JobSatisfaction' , y = 'EnvironmentSatisfaction' , data=df)
```

Out[20]:

<Axes: xlabel='JobSatisfaction', ylabel='EnvironmentSatisfaction'>



# 5. Preprocess Data

## Analyze the data

In [21]:

```
import skimpy as sk
sk.skim(df)
```

skimpy summary									
Data Summary			Data Types						
dataframe			Values	Column Type		Count			
Number of rows			1543	float64		9			
Number of columns			16	string		6			
				bool		1			
number									
column_name			NA	NA %	mean	sd	p0	p25	p50
p75	p100	hist							
Tenure			9	0.58	7.6	4.1	0	5	7
10	27								
Salary			9	0.58	67000	8500	40000	61000	66000
2000	98000								7
JobSatisfaction			28	1.81	3.4	1.1	1	3	3
4	5								
WorkLifeBalance			28	1.81	3.8	0.48	2	3.5	4
4	5								
PerformanceRating			7	0.45	3.5	0.41	1.6	3.3	3.4
3.7	5								
TrainingHours			191	12.38	35	23	5.2	18	30
40	96								
NumProjects			99	6.42	3.5	0.71	1.5	3	3.6
4	6.4								
YearsSincePromotion			1	0.06	1.2	2.1	0	0	0
2	16								
EnvironmentSatisfactio			28	1.81	2.5	0.69	1	2	2.8
3	5								
n									
bool									
column_name			Data Types						
dataframe		Values	Column Type		Count				

column_name	hist	NA	NA %	true	true rate
OverTime	1543	1			
string					
column_name	total words	NA	NA %	words per row	
Branch	1991	8	0.52	1.3	
Department	2441	0	0	1.6	
CommuteDistance	1543	0	0	1	
MaritalStatus	1543	0	0	1	
Education	2016	0	0	1.3	
ChurnLikelihood	6172	0	0	4	
End					

## Handling rows without Churned

```
In [22]:
df.dropna(subset='ChurnLikelihood', inplace=True)

print(df.info())
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Branch                                1535 non-null   object
1   Tenure                                1534 non-null   float64
2   Salary                                1534 non-null   float64
3   Department                            1543 non-null   object
4   JobSatisfaction                       1515 non-null   float64
5   WorkLifeBalance                       1515 non-null   float64
6   CommuteDistance                       1543 non-null   object
7   MaritalStatus                         1543 non-null   object
8   Education                             1543 non-null   object
9   PerformanceRating                     1536 non-null   float64
10  TrainingHours                         1352 non-null   float64
11  OverTime                              1443 non-null   object
12  NumProjects                           1444 non-null   float64
13  YearsSincePromotion                   1542 non-null   float64
14  EnvironmentSatisfaction               1515 non-null   float64
15  ChurnLikelihood                       1543 non-null   object
dtypes: float64(9), object(7)
memory usage: 193.0+ KB
None
```

Out [22]:

	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus	Education	P
0	San Francisco	4.0	63000.0	Legal	3.0	3.0	Long	Married	High School	
1	Chicago	14.0	72000.0	Accounting	4.0	4.0	Short	Single	Bachelor	
2	Miami	4.0	40000.0	Quality Assurance	3.0	3.0	Medium	Single	High School	
3	Scranton	2.0	55000.0	Legal	3.0	3.5	Short	Married	Bachelor	
4	Scranton	10.0	55500.0	Legal	3.0	3.0	Medium	Married	Bachelor	

## Handling duplicate values

In [23]:

```
print(df.duplicated().value_counts())
df.drop_duplicates()
```

False 1543
Name: count, dtype: int64

Out [23]:

	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus	Education	P
0	San Francisco	4.0	63000.0	Legal	3.0	3.000000	Long	Married	High School	
1	Chicago	14.0	72000.0	Accounting	4.0	4.000000	Short	Single	Bachelor	
2	Miami	4.0	40000.0	Quality Assurance	3.0	3.000000	Medium	Single	High School	
3	Scranton	2.0	55000.0	Legal	3.0	3.500000	Short	Married	Bachelor	
4	Scranton	10.0	55500.0	Legal	3.0	3.000000	Medium	Married	Bachelor	
...	...	...	...	...	...	...	...	...	...	...
1538	Miami	6.0	58000.0	Research and Development	4.0	4.000000	Long	Married	Bachelor	
1539	Boston	1.0	51500.0	Accounting	2.0	3.500000	Medium	Divorced	High School	
1540	San Francisco	11.0	77500.0	Human Resources	4.0	4.029174	Long	Married	Bachelor	
1541	New York	NaN	81000.0	Research and Development	4.0	3.000000	Medium	Married	Master's	
1542	Dallas	10.0	54500.0	Facilities Management	4.0	5.000000	Short	Married	High School	

1543 rows x 16 columns

## Handling high and low cardinality features

In [24]:

```
# checking the cardinality of features
df.select_dtypes("object").nunique()
```

Out[24]:

```
Branch          12
Department      15
CommuteDistance  3
MaritalStatus   3
Education       3
OverTime        1
ChurnLikelihood  3
dtype: int64
```

## Drop high cardinality features

In [25]:

```
high_cardinality = ['OverTime']

df.drop(columns=high_cardinality, inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 15 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Branch                               1535 non-null   object
 1   Tenure                               1534 non-null   float64
 2   Salary                               1534 non-null   float64
 3   Department                           1543 non-null   object
 4   JobSatisfaction                       1515 non-null   float64
 5   WorkLifeBalance                       1515 non-null   float64
 6   CommuteDistance                       1543 non-null   object
 7   MaritalStatus                         1543 non-null   object
 8   Education                             1543 non-null   object
 9   PerformanceRating                     1536 non-null   float64
10   TrainingHours                         1352 non-null   float64
11   NumProjects                           1444 non-null   float64
12   YearsSincePromotion                   1542 non-null   float64
13   EnvironmentSatisfaction               1515 non-null   float64
14   ChurnLikelihood                       1543 non-null   object
dtypes: float64(9), object(6)
memory usage: 180.9+ KB
```

## Handling empty cells

In [26]:

```
print(df.isnull().sum())
```

```
Branch          8
Tenure          9
Salary          9
Department      0
JobSatisfaction 28
WorkLifeBalance 28
CommuteDistance 0
MaritalStatus   0
Education       0
PerformanceRating 7
TrainingHours   191
NumProjects     99
YearsSincePromotion 1
EnvironmentSatisfaction 28
ChurnLikelihood 0
dtype: int64
```

In [27]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Branch                1535 non-null   object
 1   Tenure                1534 non-null   float64
 2   Salary               1534 non-null   float64
 3   Department           1543 non-null   object
 4   JobSatisfaction      1515 non-null   float64
 5   WorkLifeBalance      1515 non-null   float64
 6   CommuteDistance      1543 non-null   object
 7   MaritalStatus        1543 non-null   object
 8   Education            1543 non-null   object
 9   PerformanceRating    1536 non-null   float64
10   TrainingHours        1352 non-null   float64
11   NumProjects          1444 non-null   float64
12   YearsSincePromotion  1542 non-null   float64
13   EnvironmentSatisfaction 1515 non-null   float64
14   ChurnLikelihood      1543 non-null   object
dtypes: float64(9), object(6)
memory usage: 180.9+ KB
```

## Fill categorical value(s)

In [28]:

```
for col in df.select_dtypes(include=['object']).columns:
    df[col]=df[col].fillna(df[col].mode()[0])
```

```
df.isnull().sum()
```

Out[28]:

```
Branch                0
Tenure                9
Salary               9
Department           0
JobSatisfaction      28
WorkLifeBalance      28
CommuteDistance      0
MaritalStatus        0
Education            0
PerformanceRating    7
TrainingHours        191
NumProjects          99
YearsSincePromotion  1
EnvironmentSatisfaction 28
ChurnLikelihood      0
dtype: int64
```

In [29]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Branch                1543 non-null   object
 1   Tenure                1534 non-null   float64
 2   Salary               1534 non-null   float64
 3   Department           1543 non-null   object
 4   JobSatisfaction      1515 non-null   float64
 5   WorkLifeBalance      1515 non-null   float64
 6   CommuteDistance      1543 non-null   object
```

```

7   MaritalStatus      1543 non-null object
8   Education          1543 non-null object
9   PerformanceRating   1536 non-null float64
10  TrainingHours       1352 non-null float64
11  NumProjects         1444 non-null float64
12  YearsSincePromotion 1542 non-null float64
13  EnvironmentSatisfaction 1515 non-null float64
14  ChurnLikelihood     1543 non-null object
dtypes: float64(9), object(6)
memory usage: 180.9+ KB

```

## Fill numerical value(s)

In [30]:

```

for col in df.select_dtypes(include=['int64', 'float64']).columns:
    df[col].fillna(df[col].mean(), inplace=True)

df.isnull().sum()

```

Out[30]:

```

Branch      0
Tenure      0
Salary      0
Department  0
JobSatisfaction  0
WorkLifeBalance  0
CommuteDistance  0
MaritalStatus  0
Education   0
PerformanceRating  0
TrainingHours  0
NumProjects  0
YearsSincePromotion  0
EnvironmentSatisfaction  0
ChurnLikelihood  0
dtype: int64

```

In [31]:

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1543 entries, 0 to 1542
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Branch                1543 non-null  object
1   Tenure                1543 non-null  float64
2   Salary                1543 non-null  float64
3   Department            1543 non-null  object
4   JobSatisfaction        1543 non-null  float64
5   WorkLifeBalance        1543 non-null  float64
6   CommuteDistance        1543 non-null  object
7   MaritalStatus          1543 non-null  object
8   Education              1543 non-null  object
9   PerformanceRating      1543 non-null  float64
10  TrainingHours          1543 non-null  float64
11  NumProjects            1543 non-null  float64
12  YearsSincePromotion    1543 non-null  float64
13  EnvironmentSatisfaction 1543 non-null  float64
14  ChurnLikelihood        1543 non-null  object
dtypes: float64(9), object(6)
memory usage: 180.9+ KB

```

## Outlier Treatment



## Visualize the outliers of the numerical features

In [32]:

```
# Plot / Visualize the outliers of the numerical features
import matplotlib.pyplot as plt
import plotly.express as px

for col in df.select_dtypes(include=['int64', 'float64']).columns:
    fig = px.box(
        data_frame=df,
        x=col,
        orientation='h',
        title=f'Boxplot of the Target ({col}) - With Outliers'
    )
    fig.show()
```

## Handling Outliers

### Create a mask to filter out the outliers for 'Age'

In [33]:

```
mask_Tenure = df['Tenure'] <= 17

fig = px.box(
    data_frame=df[mask_Tenure],
    x='Tenure',
    orientation='h',
    title='Boxplot of the Target (Tenure) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'Salary'

In [34]:

```
mask_Salary1 = df['Salary'] >= 44_500
mask_Salary2 = df['Salary'] <= 88_000

fig = px.box(
    data_frame=df[mask_Salary1 & mask_Salary2],
    x='Salary',
    orientation='h',
    title='Boxplot of the Target (Salary) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'JobSatisfaction'

In [35]:

```
mask_JobSatisfaction = df['JobSatisfaction'] >= 2

fig = px.box(
    data_frame=df[mask_JobSatisfaction],
    x='JobSatisfaction',
    orientation='h',
    title='Boxplot of the Target (JobSatisfaction) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'WorkLifeBalance'

In [36]:

```
mask_WorkLifeBalance1 = df['WorkLifeBalance'] >= 2.840207
mask_WorkLifeBalance2 = df['WorkLifeBalance'] <= 4.712464

fig = px.box(
    data_frame=df[mask_WorkLifeBalance1 & mask_WorkLifeBalance2],
    x='WorkLifeBalance',
    orientation='h',
    title='Boxplot of the Target (WorkLifeBalance) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'PerformanceRating'

In [37]:

```
mask_PerformanceRating1 = df['PerformanceRating'] >= 2.835084
mask_PerformanceRating2 = df['PerformanceRating'] <= 4.13773

fig = px.box(
    data_frame=df[mask_PerformanceRating1 & mask_PerformanceRating2],
    x='PerformanceRating',
    orientation='h',
    title='Boxplot of the Target (PerformanceRating) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'TrainingHours'

In [38]:

```
mask_TrainingHours = df['TrainingHours'] <= 57

fig = px.box(
    data_frame=df[mask_TrainingHours],
    x='TrainingHours',
    orientation='h',
    title='Boxplot of the Target (TrainingHours) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'NumProjects'

In [39]:

```
mask_NumProjects = df['NumProjects'] <= 5.4

fig = px.box(
    data_frame=df[mask_NumProjects],
    x='NumProjects',
    orientation='h',
    title='Boxplot of the Target (NumProjects) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

### Create a mask to filter out the outliers for 'YearsSincePromotion'

In [40]:

```
mask_YearsSincePromotion = df['YearsSincePromotion'] <= 2
```

```
fig = px.box(
    data_frame=df[mask_YearsSincePromotion],
    x='YearsSincePromotion',
    orientation='h',
    title='Boxplot of the Target (YearsSincePromotion) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

## Create a mask to filter out the outliers for 'EnvironmentSatisfaction'

In [41]:

```
mask_EnvironmentSatisfaction = df['EnvironmentSatisfaction'] <= 4.2

fig = px.box(
    data_frame=df[mask_EnvironmentSatisfaction],
    x='EnvironmentSatisfaction',
    orientation='h',
    title='Boxplot of the Target (EnvironmentSatisfaction) - Without Outliers')

fig.update_layout(xaxis_title='Target')
fig.show()
```

## Filtering out the outliers

In [42]:

```
df = df[mask_Tenure & mask_Salary1 & mask_Salary2 & mask_JobSatisfaction & mask_WorkLife
Balance1 & mask_WorkLifeBalance2 & mask_PerformanceRating1 & mask_PerformanceRating2 & ma
sk_TrainingHours & mask_NumProjects & mask_YearsSincePromotion & mask_EnvironmentSatisfac
tion]


print(df.info())
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 904 entries, 1 to 1541
Data columns (total 15 columns):
#   Column                               Non-Null Count  Dtype
---  -
0   Branch                               904 non-null    object
1   Tenure                               904 non-null    float64
2   Salary                               904 non-null    float64
3   Department                           904 non-null    object
4   JobSatisfaction                       904 non-null    float64
5   WorkLifeBalance                       904 non-null    float64
6   CommuteDistance                       904 non-null    object
7   MaritalStatus                         904 non-null    object
8   Education                             904 non-null    object
9   PerformanceRating                     904 non-null    float64
10  TrainingHours                         904 non-null    float64
11  NumProjects                           904 non-null    float64
12  YearsSincePromotion                   904 non-null    float64
13  EnvironmentSatisfaction                904 non-null    float64
14  ChurnLikelihood                       904 non-null    object
dtypes: float64(9), object(6)
memory usage: 113.0+ KB
None
```

Out[42]:

	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus	Education
1	Chicago	14.0	72000.0	Accounting	4.0	4.0	Short	Single	Bachelor
3	Scranton	2.0	55000.0	Legal	3.0	3.5	Short	Married	Bachelor

	Branch	Tenure	Salary	Department	JobSatisfaction	WorkLifeBalance	CommuteDistance	MaritalStatus	Education
4	Scranton	10.0	55500.0	Legal	3.0	3.0	Medium	Married	Bachelor
6	Boston	10.0	82000.0	Sales	5.0	3.0	Medium	Married	Bachelor
7	New York	6.0	59000.0	Administration	3.0	3.5	Short	Divorced	High School



## Scaling

In [43]:

```
df['ChurnLikelihood'] = df['ChurnLikelihood'].map({'Slightly Likely to Churn' : 0, 'Moderately Likely to Churn' : 1, 'Highly Likely to Churn' : 2}).astype('int')
```

## 6. Modeling

### Split Data

In [44]:

```
from sklearn.model_selection import train_test_split

X = df.drop(columns=['ChurnLikelihood'], inplace=False)
y = df['ChurnLikelihood']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

In [45]:

```
# print(X)
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 904 entries, 1 to 1541
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Branch                                904 non-null    object
1   Tenure                                904 non-null    float64
2   Salary                                904 non-null    float64
3   Department                            904 non-null    object
4   JobSatisfaction                       904 non-null    float64
5   WorkLifeBalance                       904 non-null    float64
6   CommuteDistance                       904 non-null    object
7   MaritalStatus                         904 non-null    object
8   Education                             904 non-null    object
9   PerformanceRating                     904 non-null    float64
10  TrainingHours                         904 non-null    float64
11  NumProjects                           904 non-null    float64
12  YearsSincePromotion                   904 non-null    float64
13  EnvironmentSatisfaction                904 non-null    float64
dtypes: float64(9), object(5)
memory usage: 105.9+ KB
```

## Build Model

In [46]:

```
from sklearn.pipeline import make_pipeline
from category_encoders import OneHotEncoder
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression

model_dt = make_pipeline(
    OneHotEncoder(use_cat_names=True), # encode cat features
    StandardScaler(), # imputation
    DecisionTreeClassifier()) # build model


model_rf = make_pipeline(
    OneHotEncoder(use_cat_names=True), # encode cat features
    StandardScaler(), # imputation
    RandomForestClassifier()) # build model

model_lr = make_pipeline(
    OneHotEncoder(use_cat_names=True), # encode cat features
    StandardScaler(), # imputation
    LogisticRegression()) # build model

# fit the model
model_dt.fit(X_train, y_train)
model_rf.fit(X_train, y_train)
model_lr.fit(X_train, y_train)

```

Out[46]:



## Accuracy of Model

In [47]:

```

from sklearn.metrics import accuracy_score

y_pred = model_dt.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Decision Tree Accuracy:", (accuracy*100).__round__(4))

y_pred = model_rf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Random Forest Classifier:", (accuracy*100).__round__(4))

y_pred = model_lr.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Logistic Regression:", (accuracy*100).__round__(4))

```

Decision Tree Accuracy: 97.7901  
Random Forest Classifier: 98.3425  
Logistic Regression: 98.895

## Evaluate

In [48]:

```

from sklearn.metrics import mean_absolute_error

# Decision Tree
# Predict the train data
y_pred_training = model_dt.predict(X_train)
y_pred_test = model_dt.predict(X_test)

```

```

# Compute MAE
print("Training MAE:", round(mean_absolute_error(y_train, y_pred_training),2))
print("Test data MAE:", round(mean_absolute_error(y_test, y_pred_test),2))

# Random Forest Regressor
# Predict the train data
y_pred_training = model_rf.predict(X_train)
y_pred_test = model_rf.predict(X_test)

# Compute MAE
print("Training MAE:", round(mean_absolute_error(y_train, y_pred_training),2))
print("Test data MAE:", round(mean_absolute_error(y_test, y_pred_test),2))

# Logistic Regression
# Predict the train data
y_pred_training = model_lr.predict(X_train)
y_pred_test = model_lr.predict(X_test)

# Compute MAE
print("Training MAE:", round(mean_absolute_error(y_train, y_pred_training),2))
print("Test data MAE:", round(mean_absolute_error(y_test, y_pred_test),2))

```

```

Training MAE: 0.0
Test data MAE: 0.02
Training MAE: 0.0
Test data MAE: 0.02
Training MAE: 0.01
Test data MAE: 0.01

```

**The Logistic Regression Model will be used, seeing as it's accuracy is the highest and the MAE of the test data is lower/the same as the training MAE.**

## Feature Importance

### Extract the feature names and importance from the model

In [49]:

```

features = model_lr.named_steps["onehotencoder"].get_feature_names()
coefs = model_lr.named_steps["logisticregression"].coef_[0]

```

C:\Users\hroux\AppData\Roaming\Python\Python311\site-packages\category\_encoders\utils.py:366: FutureWarning:

`get\_feature\_names` is deprecated in all of sklearn. Use `get\_feature\_names\_out` instead.

In [50]:

```

import numpy as np

odds_ratios = pd.Series(np.exp(coefs), index=features).sort_values()
odds_ratios.head()

```

Out[50]:

```

Branch_Miami          0.519009
TrainingHours         0.596666
Branch_Atlanta        0.610408
Branch_Seattle        0.663997
Education_High School 0.671927
dtype: float64

```

### Plot the Feature Importance

In [51]:

```
plt.figure(figsize=(15,10))
fig = px.bar(
    data_frame=odds_ratios,
    x=odds_ratios[:10].values,
    y=odds_ratios[:10].index,
    title="Customer Churn Logistic Regression, Feature Importance (Odds Ratio)"
)

fig.update_layout(xaxis_title='Odds Ratio', yaxis_title='')
fig.show()
```

## Predict

In [52]:

```
def make_prediction(row):

    data = {
        "Branch": row[1],
        "Tenure": row[2],
        "Salary": row[3],
        "Department": row[4],
        "JobSatisfaction": row[5],
        "WorkLifeBalance": row[6],
        "CommuteDistance": row[7],
        "MaritalStatus": row[8],
        "Education": row[9],
        "PerformanceRating": row[10],
        "TrainingHours": row[11],
        "NumProjects": row[12],
        "YearsSincePromotion": row[13],
        "EnvironmentSatisfaction": row[14]
    }

    df_predict = pd.DataFrame(data, index=[0])
    prediction_dt = model_dt.predict(df_predict)[0]
    prediction_rf = model_rf.predict(df_predict)[0]
    prediction_lr = model_lr.predict(df_predict)[0]

    return f"Decision Tree: {prediction_dt}; Random Forest: {prediction_rf} Logistic Regression: {prediction_lr}"
```

In [53]:

```
print(make_prediction(pd.Series([1, 'San Francisco', 4.0, 63000.0, 'Legal', 3.0, 3.0, 'Long', 'Married', 'High School', 3.0, 88.0, 3.0, 0.0, 2.0]))) # Should output 2
print(make_prediction(pd.Series([2, 'Chicago', 14.0, 72000.0, 'Accounting', 4.0, 4.0, 'Short', 'Single', 'Bachelor', 3.6666666666666665, 30.0, 3.0, 2.0, 3.0]))) # Should output 1
```

Decision Tree: 2; Random Forest: 2 Logistic Regression: 2  
Decision Tree: 1; Random Forest: 1 Logistic Regression: 1

## Save Model

In [54]:

```
import joblib
# Save Model
joblib.dump(model_dt, '../..artifacts/model.pkl')
```

Out[54]:

```
['../..artifacts/model.pkl']
```

## Dash App

In [55]:

```

import dash
from dash import dcc, html
from dash.dependencies import Input, Output
import joblib
import pandas as pd

# Load the trained model
model = joblib.load("../artifacts/model.pkl")

# Initialize the Dash app
app = dash.Dash(__name__, external_stylesheets=['../style.css'])

# Define the layout of the app
app.layout = html.Div([
    html.H1("Employee Churn Prediction", style={'color': 'blue', 'font-size': '24px'}),

    html.Label("Branch"),
    dcc.Dropdown(
        id="branch",
        options=[
            {'label': 'San Francisco', 'value': 'San Francisco'},
            {'label': 'Chicago', 'value': 'Chicago'},
            {'label': 'Miami', 'value': 'Miami'},
            {'label': 'Scranton', 'value': 'Scranton'},
            {'label': 'Boston', 'value': 'Boston'},
            {'label': 'New York', 'value': 'New York'},
            {'label': 'Philadelphia', 'value': 'Philadelphia'},
            {'label': 'Los Angeles', 'value': 'Los Angeles'},
            {'label': 'Seattle', 'value': 'Seattle'},
            {'label': 'Atlanta', 'value': 'Atlanta'},
            {'label': 'Denver', 'value': 'Denver'},
            {'label': 'Dallas', 'value': 'Dallas'}
        ],
        value='San Francisco'
    ),

    html.Label("Tenure"),
    dcc.Input(id="tenure", type="number", value=4.0),

    html.Label("Salary"),
    dcc.Input(id="salary", type="number", value=63000.0),

    html.Label("Department"),
    dcc.Dropdown(
        id="department",
        options=[
            {'label': 'Legal', 'value': 'Legal'},
            {'label': 'Accounting', 'value': 'Accounting'},
            {'label': 'Quality Assurance', 'value': 'Quality Assurance'},
            {'label': 'Customer Service', 'value': 'Customer Service'},
            {'label': 'Sales', 'value': 'Sales'},
            {'label': 'Administration', 'value': 'Administration'},
            {'label': 'Facilities Management', 'value': 'Facilities Management'},
            {'label': 'Research and Development', 'value': 'Research and Development'},
            {'label': 'Operations', 'value': 'Operations'},
            {'label': 'Marketing', 'value': 'Marketing'},
            {'label': 'Public Relations', 'value': 'Public Relations'},
            {'label': 'IT Support', 'value': 'IT Support'},
            {'label': 'Procurement', 'value': 'Procurement'},
            {'label': 'Product Management', 'value': 'Product Management'},
            {'label': 'Human Resources', 'value': 'Human Resources'}
        ],
        value='Legal'
    ),

    html.Label("JobSatisfaction"),
    dcc.Input(id="job", type="number", value=3.0),

    html.Label("WorkLifeBalance"),
    dcc.Input(id="balance", type="number", value=3.0),

```



```

html.Label("CommuteDistance"),
dcc.Dropdown(
    id="commute",
    options=[
        {'label': 'Short', 'value': 'Short'},
        {'label': 'Medium', 'value': 'Medium'},
        {'label': 'Long', 'value': 'Long'}
    ],
    value='Long'
),

html.Label("MaritalStatus"),
dcc.Dropdown(
    id="married",
    options=[
        {'label': 'Single', 'value': 'Single'},
        {'label': 'Married', 'value': 'Married'}
    ],
    value='Married'
),

html.Label("Education"),
dcc.Dropdown(
    id="education",
    options=[
        {'label': 'High School', 'value': 'High School'},
        {'label': 'Bachelor', 'value': 'Bachelor'},
        {'label': 'Master', 'value': 'Master'},
        {'label': 'Doctor', 'value': 'Doctor'}
    ],
    value='High School'
),

html.Label("PerformanceRating"),
dcc.Input(id="rating", type="number", value=3.0),

html.Label("TrainingHours"),
dcc.Input(id="training", type="number", value=88.0),

html.Label("NumProjects"),
dcc.Input(id="projects", type="number", value=3.0),

html.Label("YearsSincePromotion"),
dcc.Input(id="promotion", type="number", value=0.0),

html.Label("EnvironmentSatisfaction"),
dcc.Input(id="satisfaction", type="number", value=2.0),

html.Button("Predict", id="predict_button", n_clicks=0),
html.Div(id="prediction_output", style={'color': 'red'})
], style={'display': 'flex', 'flexDirection': 'column', 'gap': '1em'})

```

*# Define callback to update prediction result*

```

@app.callback(
    Output("prediction_output", "children"),
    [Input("predict_button", "n_clicks")],

    [Input("branch", "value"),
     Input("tenure", "value"),
     Input("salary", "value"),
     Input("department", "value"),
     Input("job", "value"),
     Input("balance", "value"),
     Input("commute", "value"),
     Input("married", "value"),
     Input("education", "value"),
     Input("rating", "value"),
     Input("training", "value"),
     Input("projects", "value"),
     Input("promotion", "value"),
     Input("satisfaction", "value")]
)

```

```

def update_prediction(n_clicks, branch, tenure, salary, department, job, balance, commut
e, married, education, rating, training, projects, promotion, satisfaction):
    if n_clicks > 0:
        # Preprocess input data
        data = pd.DataFrame({
            "Branch": [branch],
            "Tenure": [tenure],
            "Salary": [salary],
            "Department": [department],
            "JobSatisfaction": [job],
            "WorkLifeBalance": [balance],
            "CommuteDistance": [commute],
            "MaritalStatus": [married],
            "Education": [education],
            "PerformanceRating": [rating],
            "TrainingHours": [training],
            "NumProjects": [projects],
            "YearsSincePromotion": [promotion],
            "EnvironmentSatisfaction": [satisfaction]
        })
        # Make prediction
        prediction = model.predict(data)[0]

        if prediction == 2:
            prediction_str = 'Highly Likely to Churn'
        elif prediction == 1:
            prediction_str = 'Moderately Likely to Churn'
        else:
            prediction_str = 'Slightly Likely to Churn'

        return html.Div(f"Churn Likelihood: {prediction_str}")
    else:
        return ""

# Run the app
if __name__ == "__main__":
    app.run_server(debug=True)

```

